

Fundamentals of Computers and Computational Thinking

Course Code: CS101 Credits: 4 (Lecture: 3, Practical: 2) Total Hours: 75 Hrs

MODULE I: History and Evolution of Computing System

Comprehensive Academic Notes — Units 1 through 5 (Duration: 9 Hours, Weightage: 15 Marks)

Unit 1: Evolution of Computers — History & Generations

CONCEPTUAL DEFINITION: THE MODERN COMPUTER

A **computer** is a multi-purpose, programmable, electronic data-processing machine engineered to automatically accept raw, unstructured data as input, apply high-speed mathematical, arithmetic, and logical operations in accordance with a sequence of internally stored instructions, produce structured, meaningful output information, and persist intermediate and final results in non-volatile storage media for future retrieval and processing.

1.1 The Pre-Electronic and Mechanical Computing Eras

The intellectual and technological evolution of computation spans several millennia. Human civilization has continuously sought to automate tedious mathematical calculations, reduce human error in numerical tabulation, and expand the horizons of abstract problem-solving.

The earliest known calculation aid is the **Abacus**, dating back to circa 2500 BCE in ancient Mesopotamia and subsequently perfected in China as the *Suanpan* and in Japan as the *Soroban*. The abacus utilizes a rectangular wooden frame with parallel vertical rods upon which sliding beads represent positional numerical quantities (units, tens, hundreds, thousands). Skilled operators could execute addition, subtraction, multiplication, division, and even square-root extractions at speeds comparable to early mechanical adding machines.

In 1617, Scottish mathematician John Napier published his treatise *Rabdologiae*, presenting **Napier's Bones**. This invention consisted of a set of carved ivory, wood, or bone rods inscribed with multiplication tables. Utilizing the principle of lattice multiplication (the gelosia method), Napier's bones converted complex multi-digit multiplication and division into a series of elementary single-digit additions and subtractions. Furthermore, Napier's contemporaneous discovery of logarithms enabled the creation of the **Slide Rule** (developed by William Oughtred in 1622), an analog computational tool featuring sliding logarithmic scales that remained the premier calculation instrument for scientists and engineers through the mid-20th century.

The transition from manual calculation aids to mechanical gear-driven machinery began in 1642 with the French mathematician, physicist, and philosopher Blaise Pascal. To assist his father, a regional tax administrator in Rouen, the 19-year-old Pascal designed and constructed the **Pascaline** (or *Machine Arithmétique*). The Pascaline was the first operational mechanical calculator based on a series of interlinked ten-toothed wheels representing decimal digits (**0** through **9**). A sophisticated internal gravity-assisted mechanical pawl system (the *sautoir*) executed an automatic "carry" operation whenever a wheel completed a full **360°** rotation from digit **9** to **0**. While the Pascaline could only perform addition directly and subtraction through nines' complement addition, it proved that mathematical logic could be physically embodied in mechanical clockwork.

In 1673, German polymath Gottfried Wilhelm von Leibniz dramatically expanded Pascal's concepts by inventing the **Stepped Reckoner** (*Staffelwalze*). Leibniz introduced the stepped gear cylinder (a cylindrical drum with nine teeth of incrementally increasing lengths), which allowed a movable carriage to execute direct, automated multiplication and division through repeated stepped rotations, alongside addition and subtraction. Remarkably, Leibniz was also the first Western mathematician to comprehensively document the mathematical foundations of the **Binary Number System** (using only **0** and **1**), recognizing its profound simplicity and suitability for mechanical and philosophical computation.

In 1804, a pivotal conceptual breakthrough occurred in textile manufacturing that would fundamentally alter computing history: Joseph Marie Jacquard invented the **Jacquard Loom**. Jacquard automated the weaving of intricate pictorial tapestries and damask patterns by reading sequences of rigid punched pasteboard cards. Holes punched in specific positions allowed mechanical needles to pass through and lift corresponding warp threads, while unpunched card areas blocked the needles. The Jacquard Loom introduced two revolutionary concepts: **binary-encoded data storage** (hole = 1, no hole = 0) and the principle of **program-controlled sequential machine execution** without altering physical mechanical gearing.

1.2 The Five Generations of Electronic Computing

The progression of electronic digital computers is historically categorized into five distinct generations. Each generational epoch is characterized by a revolutionary paradigm shift in active switching electronics, primary memory media, computational throughput, operating system capabilities, and programming language abstractions.

First Generation (1940–1956): The Era of Vacuum Tubes (Thermionic Valves)

The first generation of electronic digital computers relied on **vacuum tubes** (thermionic valves invented by John Ambrose Fleming and Lee de Forest) as the fundamental active switching and amplifying element. A vacuum tube consists of an evacuated glass or metal-ceramic envelope enclosing a heated filament cathode, a control grid, and an anode plate. By applying an electrical charge to the grid, the flow of thermionically emitted electrons from cathode to anode could be modulated or switched ON and OFF at kilohertz frequencies, far exceeding mechanical relays.

- **Physical & Electrical Attributes:** First-generation systems were colossal, room-sized mainframe installations weighing between **10 and 30** metric tons and occupying thousands of square feet of floor space. They consumed massive quantities of electrical power (frequently exceeding **150 kW**) and generated tremendous heat, requiring dedicated industrial air conditioning. Because vacuum tube tungsten filaments operated at incandescent temperatures, they suffered from

extremely high failure rates (Mean Time Between Failures — MTBF — was often measured in minutes or hours), requiring teams of technicians to manually diagnose and replace burnt-out tubes.

- **Memory Subsystem:** Primary memory technologies were primitive and severely capacity-constrained. Early systems employed **mercury acoustic delay lines** (where electrical pulses were converted into acoustic sound waves traveling through long tubes of liquid mercury, reflected, and recirculated), **Williams-Kilburn electrostatic tubes** (using cathode ray tube phosphor screens to store static charges), and high-speed rotating **magnetic drums** (spinning at **3,000 RPM** where magnetic tracks stored binary data). Memory capacities ranged from 1,000 to 4,000 words.
- **Software & Programming:** These machines operated without operating systems. All software was written directly in **Machine Language** (raw binary strings of 0s and 1s) or hardwired by manually configuring physical plugboards, patch cords, and rotary switches. Input and output were handled via 80-column punched cards and magnetic tape reels.
- **Seminal First-Generation Computers:**
 - **ENIAC (Electronic Numerical Integrator and Computer):** Completed in 1945 by John Mauchly and J. Presper Eckert at the University of Pennsylvania's Moore School of Electrical Engineering. Containing **17,468** vacuum tubes, **70,000** resistors, **10,000** capacitors, and **5 million** hand-soldered joints, ENIAC executed **5,000** additions per second — approximately **1,000** times faster than any electromechanical calculator of its day.
 - **EDVAC (Electronic Discrete Variable Automatic Computer):** Conceived in 1945 as the first binary, stored-program electronic computer.
 - **UNIVAC I (Universal Automatic Computer I):** Delivered to the United States Census Bureau in 1951, UNIVAC I was the world's first commercially mass-manufactured general-purpose electronic computer. It achieved worldwide acclaim when it accurately predicted the landslide victory of Dwight D. Eisenhower in the 1952 US Presidential Election on live national television.

Second Generation (1956–1963): The Transistor Revolution

The invention of the point-contact transistor at Bell Laboratories in December 1947 by physicists John Bardeen, Walter Brattain, and William Shockley (who shared the 1956 Nobel Prize in Physics) catalyzed the second generation of computing. The transistor is a three-terminal solid-state semiconductor device that controls current flow through solid silicon or germanium crystals without requiring heated filaments or a vacuum.

- **Technological Advantages:** Discrete transistors reduced physical component volume by a factor of **1,000**, slashed electrical power consumption by over **95%**, generated minimal heat, operated at megahertz switching frequencies, and eliminated filament burnout, boosting machine reliability by multiple orders of magnitude. Computational execution times dropped from milliseconds (10^{-3} s) to microseconds (10^{-6} s).
- **Memory Revolution — Magnetic Core Memory:** Second-generation systems replaced acoustic delay lines with **Magnetic Core Memory**. Developed by Jay Forrester and An Wang, core memory comprised tiny toroidal ferrite rings (cores) threaded onto a two-dimensional orthogonal wire grid (X and Y drive lines, plus sense and inhibit wires). By passing half-select currents along intersecting wires, individual cores were magnetized clockwise (representing binary **1**) or counter-clockwise (representing binary **0**). Core memory provided non-volatile, random-access memory (RAM) with access cycle times of **1 to 2 microseconds**.
- **Software Advancements:** The second generation saw the birth of **Assembly Languages** (which replaced machine code with mnemonic abbreviations such as **MOV**, **ADD**, **SUB**, and symbolic memory addresses) and the emergence of the first standardized **High-Level Programming Languages (HLL)**:
 - **FORTRAN (Formula Translation):** Developed in 1957 by John Backus at IBM for scientific and engineering calculation.
 - **COBOL (Common Business-Oriented Language):** Formulated in 1959 by the CODASYL committee, heavily guided by Rear Admiral Grace Hopper, designed for human-readable business data processing.
 - **ALGOL (Algorithmic Language):** Introduced in 1958/1960, establishing block structures, lexical scoping, and formal grammar syntax (Backus-Naur Form - BNF).

Primitive resident monitor programs emerged, serving as early ancestors of batch-processing operating systems.

- **Representative Systems:** IBM 7090, IBM 1401 (the most popular commercial computer of its era, with over 12,000 units leased), CDC 1604, and UNIVAC 1107.

Third Generation (1964–1971): Integrated Circuits (ICs) & Time-Sharing

The third generation began with the commercial deployment of the **Integrated Circuit (IC)**, independently invented by Jack Kilby of Texas Instruments (who created a hybrid germanium circuit in 1958) and Robert Noyce of Fairchild Semiconductor (who patented the monolithic silicon planar IC in 1959).

- **Monolithic Integration:** Instead of soldering thousands of discrete transistors, diodes, and resistors onto circuit boards, the planar process allowed entire circuits containing dozens to hundreds of electronic components to be microscopically etched and diffused onto a single monolithic wafer of monocrystalline silicon. This era spanned **Small Scale Integration (SSI, < 10 logic gates per chip)** and **Medium Scale Integration (MSI, 10 to 100 gates per chip)**.
- **Performance and Storage:** Computational instruction cycle speeds reached nanoseconds (10^{-9} s). Semiconductor solid-state RAM chips (such as the Intel 1103 DRAM chip introduced in 1970) began displacing magnetic core memory. Mass-storage magnetic disk packs with movable read/write heads provided tens of megabytes of online, random-access storage.
- **Operating Systems & Multiprogramming:** The third generation revolutionized software architecture. Instead of processing single batch jobs sequentially, sophisticated **Multiprogramming and Time-Sharing Operating Systems** (such as IBM OS/360, MIT's CTSS, and MULTICS) were developed. Multiple concurrent users could interact with a central mainframe simultaneously via remote alphanumeric keyboard-and-monitor video display terminals (VDTs).
- **Seminal Architectures:**
 - **IBM System/360 (1964):** A monumental **\$5 billion** engineering effort that unified disparate commercial and scientific architectures into a single scalable family of software-compatible machines, establishing the standard 8-bit byte and 32-bit word.
 - **DEC PDP-8 (1965):** Engineered by Digital Equipment Corporation, the PDP-8 was the world's first commercially triumphant **Minicomputer**, bringing computing out of specialized air-conditioned data centers directly into university laboratories and manufacturing floors at a fraction of mainframe costs.

Fourth Generation (1971–Present): Microprocessors, VLSI & The Personal Computer

The defining technological milestone of the fourth generation was the invention of the **Microprocessor** — the complete fabrication of a computer's Central Processing Unit (ALU, Control Unit, and internal registers) onto a single silicon semiconductor chip.

- **The Intel 4004 & Microprocessor Evolution:** In 1971, engineers Ted Hoff, Federico Faggin, Stan Mazor, and Masatoshi Shima designed the **Intel 4004** (a 4-bit microprocessor comprising **2,300** transistors on a **12 mm²** die running at **740 kHz**). This was rapidly followed by the 8-bit Intel 8008, the 8080 (1974), and the legendary 16-bit Intel 8086 (1978), which established the enduring x86 instruction set architecture.

- **Integration Density Scaling:** Semiconductor manufacturing progressed through **Large Scale Integration (LSI, 1,000 to 20,000 transistors)** to **Very Large Scale Integration (VLSI, 20,000 to 1,000,000+ transistors)**. **Moore's Law** — the empirical observation formulated by Gordon Moore in 1965 stating that the number of transistors on a microchip doubles approximately every 18 to 24 months while the cost per transistor halves — served as the driving roadmap for the microelectronics industry for half a century.
- **The Microcomputer & PC Revolution:** Microprocessors democratized computing. The MITS Altair 8800 (1975) sparked the home hobbyist movement, leading to Steve Wozniak and Steve Jobs creating the Apple II (1977), the Commodore PET, and the historic **IBM Personal Computer (Model 5150)** in August 1981, powered by the Intel 8088 CPU and Microsoft MS-DOS.
- **Graphical User Interfaces & Networking:** The fourth generation witnessed the transition from command-line interfaces to **Graphical User Interfaces (GUI)** with the WIMP (Windows, Icons, Menus, Pointers) paradigm (pioneered at Xerox PARC and popularized by the Apple Macintosh in 1984 and Microsoft Windows). Simultaneously, local networking (Ethernet) and the global interconnected network of networks — the **Internet** (governed by the TCP/IP protocol suite) and the World Wide Web (invented by Tim Berners-Lee in 1989) — transformed computers from isolated calculating tools into global communications appliances.

Fifth Generation (Present & Emerging Future): ULSI, AI & Quantum Paradigms

The fifth generation transcends classical transistor dimensional scaling, moving into the realm of **Ultra Large Scale Integration (ULSI)** (chips exceeding **100 billion** transistors), massively parallel multi-core architectures, specialized hardware accelerators for Artificial Intelligence, and non-von Neumann computing systems.

- **Key Pillars:** Artificial Intelligence (AI), Machine Learning (ML), Deep Neural Network training and inference, Natural Language Processing (NLP), Computer Vision, and Autonomous Robotics.
- **Specialized Silicon Accelerators:** Massive deployment of Graphics Processing Units (GPUs), Tensor Processing Units (TPUs), and Neural Processing Units (NPUs) engineered with dedicated matrix-multiplication systolic arrays.
- **Emerging Paradigms:** Quantum Processing Units (QPU) exploiting quantum superposition and entanglement, Neuromorphic Computing (silicon spiking neural networks mimicking biological synapses), and Optical/Photonic Computing utilizing photons for speed-of-light data transmission.

1.3 Comprehensive Generational Comparison Table

Generation	Active Switching Element	Primary Memory Technology	Operating Speed	Software / OS Architecture	Typical Input / Output	Seminal Computing Machines
1st Gen (1940–56)	Thermionic Vacuum Tubes / Valves	Acoustic Delay Lines, Magnetic Drums	Milliseconds (10⁻³ s) ~thousands ops/sec	Raw Machine Language, Plugboards, No OS	Punched cards, paper tape, magnetic tape	ENIAC, EDVAC, EDSAC, UNIVAC I, IBM 701
2nd Gen (1956–63)	Discrete Bipolar Transistors (BJT)	Magnetic Ferrite Core Memory	Microseconds (10⁻⁶ s) ~hundreds of thousands ops/sec	Assembly Language, FORTRAN, COBOL, Early Batch Monitors	Punched cards, high-speed magnetic tape, early disk	IBM 7094, IBM 1401, CDC 1604, UNIVAC 1107
3rd Gen (1964–71)	Integrated Circuits (SSI & MSI on silicon)	Semiconductor RAM chips (DRAM), Core	Nanoseconds (10⁻⁹ s) ~millions ops/sec (MIPS)	Time-Sharing OS, Multiprogramming (OS/360, MULTICS, UNIX)	Keyboards, Alphanumeric Video Display Terminals (VDTs)	IBM System/360, DEC PDP-8, PDP-11, CDC 6600
4th Gen (1971–Pres.)	Microprocessors (LSI, VLSI, ULSI)	Semiconductor DRAM, SRAM, NVMe Flash SSDs	Picoseconds (10⁻¹² s) ~billions ops/sec (GFLOPS/TFLOPS)	GUI Operating Systems (Windows, Linux, macOS), C++, Java, Python	Mouse, Optical Keyboards, Touchscreens, 4K Monitors	Intel Core, AMD Ryzen, Apple Silicon M-Series, IBM PC
5th Gen (Emerging)	ULSI, Quantum Gates, 3D Chiplets, NPUs	High Bandwidth Memory (HBM3), Quantum States	Massively Parallel / ExaFLOPS (10¹⁸ ops/sec)	AI Frameworks (PyTorch, TensorFlow), Qiskit, Quantum OS	Natural Voice Interfaces, Neural BCI, Computer Vision	NVIDIA H100/B200, IBM Quantum Condor, Google Sycamore

Unit 2: Overview of Computer System — Von Neumann Model & Number Systems

2.1 The Von Neumann Computer Architecture

In June 1945, the brilliant Hungarian-American mathematician and polymath John von Neumann authored a landmark report entitled *"First Draft of a Report on the EDVAC"*. In this manuscript, von Neumann articulated an architectural blueprint for a general-purpose electronic computer that separated the machine's physical hardware from its logical programming.

Prior to the von Neumann model, early machines like the ENIAC were "fixed-program" or hardwired computers; if an engineer wished to switch from calculating an artillery trajectory to computing a nuclear hydrodynamic equation, the machine had to be powered down, manually rewired with thousands of patch cords, and

physically reconfigured over several days.

The revolutionary insight of the **Stored-Program Concept** is that:

THE STORED-PROGRAM PRINCIPLE

Both program instructions (the operational algorithms) and the computational data (the operands) are encoded in identical binary format and stored together in the same shared, read-write primary memory space. The CPU sequentially fetches an instruction from memory, interprets its operation code (opcode), executes the corresponding micro-operations on data fetched from memory, and writes results back to memory using a unified bus interconnect.




Figure 1.1: Complete Von Neumann Architecture with CPU Subsystems, Main Memory, and Tri-Bus Interconnect.

Detailed Subsystems of the Von Neumann Model

- The Central Processing Unit (CPU):** The computational brain of the system, comprising:
 - The Control Unit (CU):** The master coordinator and traffic controller of the computer. The CU generates synchronization clock pulses and control signals that orchestrate the fetching, decoding, and execution of instructions. It directs the movement of data between memory, ALU, internal registers, and input/output modules.
 - The Arithmetic and Logic Unit (ALU):** The functional unit that executes all computational operations:
 - Arithmetic Operations:** Addition, subtraction, multiplication, and division on binary integers and floating-point numbers.
 - Logical Operations:** Bitwise boolean operations (AND, OR, NOT, XOR, NOR, NAND) and comparison evaluations ($<$, $\leq$, $>$, $\geq$, $==$, $\neq$).
 - Bit-Shifting Operations:** Logical left/right shift, arithmetic right shift (preserving sign bits), and cyclic rotation.
 - The Internal CPU Register File:** Ultra-fast, low-capacity storage cells fabricated directly inside the CPU using static flip-flop circuits:
 - Program Counter (PC):** Holds the exact physical memory address of the *next* sequential instruction to be fetched from memory. It automatically increments after each fetch, unless redirected by a branch, jump, or subroutine call.
 - Instruction Register (IR):** Holds the binary opcode and operand specifiers of the instruction currently undergoing decoding and execution.
 - Memory Address Register (MAR):** Holds the physical memory address destined to be driven across the Address Bus during a memory read or write cycle.
 - Memory Buffer Register (MBR) / Memory Data Register (MDR):** Holds the actual data word that has just been read from memory or is waiting to be written to memory across the Data Bus.
 - Accumulator (ACC):** A primary general-purpose working register that holds the immediate operands and intermediate mathematical results generated by the ALU.
 - Program Status Word (PSW) / Flags Register:** A collection of individual status flip-flops (Zero Flag Z , Carry Flag C , Sign/Negative Flag S , Overflow Flag V , Parity Flag P , Interrupt Enable Flag I) updated by the ALU after each mathematical operation.
- The Main Memory Subsystem (RAM):** A linear, contiguous array of addressable storage cells (bytes). Each storage cell is assigned a unique non-negative integer address ranging from 0 to $2^N - 1$, where N represents the bit-width of the memory address bus. In the von Neumann model, instructions and data reside in this single uniform address space.
- The Input/Output (I/O) Subsystem:** Interfaces the high-speed internal CPU and memory with external human and machine peripherals (keyboards, monitors, network transceivers, secondary disk drives). It performs signal voltage level translation, serial-to-parallel data conversion, protocol handshaking, and buffering.
- The System Bus Interconnection Subsystem:** The shared electrical parallel highway connecting the CPU, Memory, and I/O, divided into three specialized functional buses:
 - Address Bus (Unidirectional):** Transmits memory and I/O port addresses originating exclusively from the CPU to external memory chips and controllers. An N -bit address bus can address 2^N distinct bytes of physical memory (e.g., a 32-bit address bus addresses $2^{32} = 4$ GB; a 64-bit address bus addresses

$2^{64} = 16 \text{ Exabytes}$).

- **Data Bus (Bidirectional):** Carries bidirectional data words and instruction opcodes between the CPU, memory, and I/O devices. The width of the data bus (W bits, e.g., 64 bits) determines the system's fundamental word transfer capacity per clock cycle.
- **Control Bus (Bidirectional):** Carries command, timing, and synchronization signals across the system (e.g., Memory Read **MEMR**, Memory Write **MEMW**, I/O Read **IOR**, I/O Write **IOW**, System Clock **CLK**, Interrupt Request **INTR**, Interrupt Acknowledge **INTA**, Bus Request **BREQ**, and Ready signals).

THE VON NEUMANN BOTTLENECK & THE MEMORY WALL

Because both instruction fetches and data read/writes must share the **same physical bus interconnect** in the von Neumann architecture, the CPU cannot read an instruction and read/write operand data at the exact same clock cycle.

Over decades, CPU processing frequencies grew exponentially (**50%** per year), whereas DRAM memory access latency improved at a sluggish pace (**7%** per year). This divergence created the **Memory Wall** — the CPU spends a substantial percentage of its clock cycles idle (in "wait states") waiting for data transfers across the system bus.

Mitigations in Modern Computer Architecture:

- **Multi-Level On-Die Cache Hierarchies (L1, L2, L3):** Placing ultra-high-speed SRAM caches directly onto the CPU die to satisfy up to **95%** of memory requests within 1 to 10 clock cycles.
- **Modified Harvard Architecture:** Implementing physically separated, independent L1 Instruction Cache (L1 I-Cache) and L1 Data Cache (L1 D-Cache) inside the CPU core, allowing simultaneous instruction fetch and data access, while maintaining a unified memory model in DRAM.
- **Hardware Branch Prediction & Speculative Execution:** Predicting execution paths ahead of time and prefetching data into caches before instructions are reached.

The Formal Four-Phase Instruction Execution Cycle

The CPU executes continuous programs by repeatedly cycling through four deterministic phases:

1. Fetch Phase:

- The content of the Program Counter (PC) is transferred to the Memory Address Register: $MAR \leftarrow PC$.
- The Control Unit asserts the **MEMR** signal on the control bus.
- The memory subsystem reads the instruction located at address MAR and places it onto the Data Bus, which is latched into the Memory Data Register: $MDR \leftarrow M[MAR]$.
- The instruction is copied from MDR to the Instruction Register: $IR \leftarrow MDR$.
- The Program Counter is automatically incremented to point to the next instruction in memory: $PC \leftarrow PC + \text{instruction_size}$.

2. Decode Phase:

- The Control Unit decodes the opcode bits stored in the IR using internal instruction decoders.
- Effective memory addresses of any required operands are calculated (resolving direct, indirect, register, or indexed addressing modes).

3. Execute Phase:

- If operands reside in memory, they are fetched via $MAR \leftarrow \text{effective_address}$ and $MDR \leftarrow M[MAR]$.
- The Control Unit signals the ALU to perform the specified arithmetic or logical operation (e.g., adding MDR to the Accumulator $ACC \leftarrow ACC + MDR$).
- If the instruction is a conditional branch (e.g., **JZ** — Jump if Zero) and the Zero Flag $Z = 1$, the PC is updated: $PC \leftarrow \text{target_address}$.

4. Store / Write-Back Phase:

- The final computed result in the Accumulator or working register is written back to the destination general-purpose register or to primary memory: $MAR \leftarrow \text{dest_address}, MDR \leftarrow \text{result}, M[MAR] \leftarrow MDR$.
- The flags register (PSW) is updated to reflect negative, zero, carry, or overflow conditions.

2.2 Positional Number Systems in Digital Computing

Digital computers are electronic switching machines. Because semiconductor transistors operate with maximum noise immunity when switching between two stable saturation states (cut-off = **0V** / Logic **0**, and saturation = **+1.8V** or **+3.3V** / Logic **1**), all computer arithmetic is founded upon positional number systems.

In a positional number system with base (radix) r , a number N possessing n integer digits and m fractional digits is represented as:

$$N = (d_{n-1}d_{n-2} \dots d_1d_0.d_{-1}d_{-2} \dots d_{-m})_r$$

where each digit d_i satisfies $0 \leq d_i < r$. The precise decimal magnitude of the number is calculated via polynomial expansion:

$$V(N) = \sum_{i=-m}^{n-1} d_i \cdot r^i = d_{n-1}r^{n-1} + d_{n-2}r^{n-2} + \dots + d_1r^1 + d_0r^0 + d_{-1}r^{-1} + \dots + d_{-m}r^{-m}$$

The Four Primary Number Systems

Number System	Radix (r)	Permissible Digits / Symbols	Positional Weight Multipliers	Primary Utility in Computing
Binary	r = 2	0, 1	..., 2 ³ = 8, 2 ² = 4, 2 ¹ = 2, 2 ⁰ = 1 · 2 ⁻¹ = 0.5, 2 ⁻² = 0.25	Physical hardware storage, logic gates, ALU arithmetic
Octal	r = 8	0, 1, 2, 3, 4, 5, 6, 7	..., 8 ² = 64, 8 ¹ = 8, 8 ⁰ = 1 · 8 ⁻¹ = 0.125, 8 ⁻² = 0.015625	Unix/Linux file permissions (e.g. <code>chmod 777</code>), legacy minicomputers
Decimal	r = 10	0, 1, 2, 3, 4, 5, 6, 7, 8, 9	..., 10 ² = 100, 10 ¹ = 10, 10 ⁰ = 1 · 10 ⁻¹ = 0.1, 10 ⁻² = 0.01	Human interface, accounting, financial and business software
Hexadecimal	r = 16	0–9, A(10), B(11), C(12), D(13), E(14), F(15)	..., 16 ² = 256, 16 ¹ = 16, 16 ⁰ = 1 · 16 ⁻¹ = 0.0625	Memory address notation (<code>0x7FFF80</code>), IPv6, MAC addresses, RGB colors

Unit 3: Number Conversions & Digital Codes

3.1 Systematic Base Conversion Procedures

Translating numbers between bases is an essential mathematical operation in computer science. Integer and fractional components are converted using complementary, inverse algorithms.

1. Decimal to Any Base r

- **Integer Part:** Apply repeated successive division by target radix **r**. In each step, record the integer quotient and the remainder ($0 \leq \text{remainder} < r$). Repeat until the quotient reaches **0**. The converted integer is formed by collecting the remainders in **reverse chronological order (from bottom to top, representing Most Significant Digit to Least Significant Digit)**.
- **Fractional Part:** Apply repeated successive multiplication by target radix **r**. In each step, record the resulting integer digit and retain the new fractional remainder for subsequent multiplication. Terminate when the fraction reaches exact **0.0** or when the required degree of precision (number of fractional digits) is attained. The converted fraction is formed by collecting the integer digits in **direct chronological order (from top to bottom)**.

WORKED EXAMPLE 3.1: DECIMAL TO BINARY & HEXADECIMAL

Problem: Convert decimal number $(109.6875)_{10}$ into Binary and Hexadecimal.

Step 1: Decimal to Binary

Integer Part $(109)_{10}$:

- $109 \div 2 = 54$ with remainder 1 (LSD b_0)
- $54 \div 2 = 27$ with remainder 0
- $27 \div 2 = 13$ with remainder 1
- $13 \div 2 = 6$ with remainder 1
- $6 \div 2 = 3$ with remainder 0
- $3 \div 2 = 1$ with remainder 1
- $1 \div 2 = 0$ with remainder 1 (MSD b_6)

Reading bottom-to-top: $(109)_{10} = (1101101)_2$.

Fractional Part $(0.6875)_{10}$:

- $0.6875 \times 2 = 1.3750 \Rightarrow$ Record 1, fraction remains 0.3750
- $0.3750 \times 2 = 0.7500 \Rightarrow$ Record 0, fraction remains 0.7500
- $0.7500 \times 2 = 1.5000 \Rightarrow$ Record 1, fraction remains 0.5000
- $0.5000 \times 2 = 1.0000 \Rightarrow$ Record 1, fraction is 0.0000 (Done)

Reading top-to-bottom: $(0.6875)_{10} = (0.1011)_2$.

Combined Binary Result: $(109.6875)_{10} = (1101101.1011)_2$.

Step 2: Decimal to Hexadecimal

Integer Part $(109)_{10}$:

- $109 \div 16 = 6$ remainder 13 $\Rightarrow$ D
- $6 \div 16 = 0$ remainder 6

Reading bottom-to-top: $(109)_{10} = (6D)_{16}$.

Fractional Part $(0.6875)_{10}$:

- $0.6875 \times 16 = 11.0000 \Rightarrow$ Record 11 $\Rightarrow$ B (Done)

Combined Hexadecimal Result: $(109.6875)_{10} = (6D.B)_{16}$.

2. Fast Direct Grouping: Binary $\leftrightarrow$ Octal $\leftrightarrow$ Hexadecimal

Because the bases are powers of 2 ($8 = 2^3$ and $16 = 2^4$), conversion between binary and octal or hexadecimal requires no arithmetic calculations — only simple bit clustering:

- **Binary to Octal:** Group bits into clusters of **3** starting from the radix point (move left for integer part, padding leading zeros; move right for fractional part, padding trailing zeros). Replace each 3-bit group with its corresponding octal digit (**0–7**).
- **Binary to Hexadecimal:** Group bits into clusters of **4** starting from the radix point (pad zeros left and right as needed). Replace each 4-bit nibble with its hexadecimal symbol (**0–9, A–F**).
- **Hexadecimal to Octal:** Convert each hexadecimal digit into its exact 4-bit binary equivalent, and then regroup the combined binary string into 3-bit octal clusters.

WORKED EXAMPLE 3.2: GROUPING CONVERSIONS

Convert $(10110110101.110101)_2$ to Octal and Hexadecimal:

To Octal (3-bit clusters):

$$\underbrace{010}_2 \underbrace{110}_6 \underbrace{110}_6 \underbrace{101}_5 \cdot \underbrace{110}_6 \underbrace{101}_5 \Rightarrow (2665.65)_8$$

To Hexadecimal (4-bit clusters):

$$\underbrace{0101}_5 \underbrace{1011}_B \underbrace{0101}_5 \cdot \underbrace{1101}_D \underbrace{0100}_4 \Rightarrow (5B5.D4)_{16}$$

3.2 Complements & Signed Binary Arithmetic

In digital processors, subtraction is executed using addition hardware by encoding negative numbers in **Complements**.

- **1's Complement:** Formed by bitwise inverting all bits ($0 \rightarrow 1, 1 \rightarrow 0$). For an n -bit number N , $1's \text{ Comp}(N) = (2^n - 1) - N$. Disadvantages: dual representations for zero ($+0 = 0000_2, -0 = 1111_2$) and requires an "end-around carry" addition stage.
- **2's Complement:** The universal standard in modern computing architectures. Formed by adding 1 to the 1's complement:

$$2's \text{ Comp}(N) = [1's \text{ Comp}(N)] + 1 = 2^n - N$$

Key Properties:

- Provides a single unique representation for zero (00000000_2).
- An n -bit 2's complement register represents signed integers in the range $-2^{n-1} \leq N \leq +2^{n-1} - 1$ (e.g., 8-bit signed range is -128 to $+127$).
- Subtractions are computed via simple binary addition: $A - B = A + [2's \text{ Comp}(B)]$. Any final carry-out bit beyond the n -th bit is simply discarded.

WORKED EXAMPLE 3.3: 2'S COMPLEMENT SUBTRACTION

Evaluate $(45)_{10} - (23)_{10}$ using 8-bit 2's complement arithmetic:

$$+45 = (00101101)_2$$

$$+23 = (00010111)_2$$

$$1's \text{ Complement of } 23 = 11101000_2$$

$$2's \text{ Complement of } 23 = 11101000_2 + 1 = (11101001)_2$$

Now execute binary addition:

$$\begin{array}{r} 00101101 \quad (+45) \\ +11101001 \quad (-23 \text{ in } 2's \text{ complement}) \\ \hline 1\ 00010110 \quad (\text{Final carry 1 is discarded}) \end{array}$$

The result is $(00010110)_2 = 16 + 4 + 2 = (22)_{10}$, which is correct ($45 - 23 = 22$).

3.3 Digital Codes

Digital codes are structured binary representations engineered for computation, data transmission, or electromechanical sensing.

1. Binary Coded Decimal (BCD / 8421 Code)

BCD represents each decimal digit (0 through 9) independently using a 4-bit binary nibble with weights **8, 4, 2, 1**.

- Decimal $(852)_{10}$ in BCD is: $\underbrace{1000}_8 \underbrace{0101}_5 \underbrace{0010}_2$.
- **Invalid BCD States:** A 4-bit nibble can represent 16 states (0000_2 to 1111_2). States 1010_2 (10_{10}) through 1111_2 (15_{10}) are **illegal/invalid BCD codes**.
- **BCD Addition Correction Rule:** When adding two BCD digits, if the sum exceeds **9** (1001_2) or generates an internal binary carry, the result contains an invalid code. To correct the error and propagate a proper decimal carry, add $(0110)_2$ (6_{10}) to that specific 4-bit sum.

2. Excess-3 Code (XS-3)

Excess-3 is an **unweighted, self-complementing** decimal code formed by adding $(0011)_2$ (3_{10}) to the 8421 BCD representation of each digit.

- **Self-Complementing Characteristic:** The 1's complement (bitwise NOT) of the Excess-3 code of digit d directly yields the Excess-3 representation of its 9's complement ($9 - d$). For instance, the Excess-3 code for **2** is 0101_2 ; inverting all bits yields 1010_2 , which is the Excess-3 code for **7** ($9 - 2 = 7$). This property greatly simplified subtraction circuitry in early computing ALUs.

3. Gray Code (Reflected Binary Code)

Gray code is an **unweighted, unit-distance** binary code in which any two successive numerical values differ in exactly **one single bit position**.

- **Application in Shaft Encoders:** When an optical or mechanical shaft angle encoder rotates between positions, standard binary transitions (such as $3 = 011_2 \rightarrow 4 = 100_2$) require multiple bits to change simultaneously. Due to microscopic mechanical tolerances, switches never change at the exact same microsecond, causing transient intermediate reading errors (e.g., momentarily reading $111_2 = 7$). Gray code guarantees that only one bit changes per angular increment, completely eliminating reading glitches.
- **Binary to Gray Conversion Algorithm:** Let Binary be $B = b_{n-1}b_{n-2} \dots b_0$ and Gray be $G = g_{n-1}g_{n-2} \dots g_0$:

$$g_{n-1} = b_{n-1} \quad (\text{MSB is preserved})$$

$$g_i = b_{i+1} \oplus b_i \quad \text{for } i = n-2, \dots, 0$$

- **Gray to Binary Conversion Algorithm:**

$$b_{n-1} = g_{n-1}$$

$$b_i = b_{i+1} \oplus g_i \quad \text{for } i = n-2, \dots, 0$$

WORKED EXAMPLE 3.4: GRAY CODE CONVERSIONS

1. Convert Binary $(1011)_2$ to Gray Code:

- $g_3 = b_3 = 1$
- $g_2 = b_3 \oplus b_2 = 1 \oplus 0 = 1$
- $g_1 = b_2 \oplus b_1 = 0 \oplus 1 = 1$
- $g_0 = b_1 \oplus b_0 = 1 \oplus 1 = 0$

Result: $(1011)_2 = (1110)_{\text{Gray}}$.

2. Convert Gray Code $(1110)_{\text{Gray}}$ back to Binary:

- $b_3 = g_3 = 1$
- $b_2 = b_3 \oplus g_2 = 1 \oplus 1 = 0$
- $b_1 = b_2 \oplus g_1 = 0 \oplus 1 = 1$
- $b_0 = b_1 \oplus g_0 = 1 \oplus 0 = 1$

Result: $(1110)_{\text{Gray}} = (1011)_2$.

Unit 4: Pioneers & Contributors of Computing Systems

The architectural and theoretical edifices of computer science are the direct culmination of seminal discoveries made by visionary thinkers. This section details the explicit contributions of the five monumental pioneers specified in the syllabus.

4.1 Charles Babbage (1791–1871) — The First Mechanical Computer

Celebrated historically as the "*Father of the Computer*", British mathematician, mechanical engineer, and polymath Charles Babbage conceived the earliest programmable computing machines during the 19th century.

- **The Difference Engine (1822):** Babbage was deeply frustrated by the pervasive human errors in printed navigation, astronomical, and logarithmic tables. He engineered the Difference Engine — a massive mechanical special-purpose calculator that computed polynomial functions automatically using the mathematical **method of finite differences**. By eliminating multiplication and division in favor of pure repeated additions, the machine could tabulate mathematical tables and automatically imprint results onto stereotype printing plates.
- **The Analytical Engine (1837):** Babbage's crowning intellectual masterpiece — the world's first design for a general-purpose, programmable, fully automatic mechanical computer. The Analytical Engine incorporated every fundamental architectural component found in modern electronic systems:
 - **The "Store":** A mechanical memory unit capable of storing 1,000 numbers of 50 decimal digits each.
 - **The "Mill":** The arithmetic processing unit containing geared columns that executed addition, subtraction, multiplication, and division under mechanical control.
 - **Control & Program Input:** Utilized strings of Jacquard punched cards to input operational instructions and variable values. The machine supported conditional branching (if-then testing of mechanical sign levers) and looping.
 - **Output Unit:** An automated typesetting mechanism and curve plotter.
- **Ada Lovelace (1815–1852):** Augusta Ada King, Countess of Lovelace, translated Luigi Menabrea's French paper on Babbage's Analytical Engine, appending extensive "Notes". Note G contained an algorithm to compute Bernoulli numbers using the engine — recognized today as the **world's first computer program**. Lovelace made the profound philosophical leap that Babbage's engine could manipulate any entities governed by rules (including music and symbolic logic), establishing the concept of universal symbolic computation.

4.2 John von Neumann (1903–1957) — The Stored-Program Architecture

Hungarian-American mathematician John von Neumann revolutionized mathematics, quantum mechanics, nuclear physics, game theory, and computer systems design.

- **First Draft of a Report on the EDVAC (1945):** Von Neumann formally synthesized the **stored-program architecture**. He established the standard model comprising a Central Processing Unit (ALU, Control Unit, Register File), a unified primary memory holding both instructions and data, and Input/Output subsystems interconnected via buses.
- **The IAS Machine (1952):** Built at the Institute for Advanced Study in Princeton, the IAS computer served as the physical architectural prototype for the entire mainframe industry (including the IBM 701) and established the foundational 40-bit word and parallel arithmetic processing standards.
- **Theoretical Advances:** Authored foundational mathematical papers on self-reproducing cellular automata, stochastic computing, and reliable computing architectures constructed from unreliable components.

4.3 Alan Turing (1912–1954) — Theoretical Foundations of Computer Science

British mathematician, logician, and wartime cryptanalyst Alan Mathison Turing is universally revered as the "*Father of Theoretical Computer Science and Artificial Intelligence*".

- **The Universal Turing Machine (1936):** In his landmark paper "*On Computable Numbers, with an Application to the Entscheidungsproblem*", Turing formulated an abstract mathematical model of computation consisting of an infinite tape divided into discrete cells, a read/write head, and a state register

governed by a finite transition table. Turing proved that a single **Universal Turing Machine (UTM)** could simulate any algorithmic computation whatsoever if encoded as a program on its tape. He proved the fundamental mathematical limits of computation, demonstrating the unsolvability of the **Halting Problem** (it is undecidable whether an arbitrary program will terminate or run forever).

- **Wartime Cryptanalysis at Bletchley Park:** During World War II, Turing was the intellectual leader at Hut 8. He designed the electromechanical **Bombe** machine, which broke the German military Enigma cipher at industrial speed, decrypting millions of military communications (Ultra intelligence) that historians estimate shortened the war in Europe by two to four years, saving millions of lives.
- **The Turing Test for Artificial Intelligence (1950):** In "*Computing Machinery and Intelligence*", Turing proposed the definitive operational benchmark for machine intelligence (the *Imitation Game*), asserting that if a human interrogator communicating via text cannot reliably tell a machine from a human, the machine exhibits true intelligent behavior.

4.4 John Mauchly (1907–1980) & J. Presper Eckert (1919–1995) — ENIAC & UNIVAC

Physicist John Mauchly and electrical engineering genius J. Presper Eckert Jr. formed the legendary hardware engineering partnership at the University of Pennsylvania's Moore School.

- **ENIAC (1943–1945):** Conceived to calculate artillery firing tables for the US Army Ballistics Research Laboratory, ENIAC was the world's **first operational, general-purpose electronic digital computer**. Eckert's engineering brilliance ensured unprecedented vacuum tube reliability by operating tube heaters at conservative voltages, proving that colossal electronic machines could function reliably.
- **EDVAC & BINAC:** Co-designed EDVAC and engineered BINAC (Binary Automatic Computer), the first stored-program computer in the US to utilize magnetic tape storage.
- **UNIVAC I & The Commercial Computer Industry:** Founded the Eckert-Mauchly Computer Corporation (EMCC). In 1951, they delivered **UNIVAC I**, marking the birth of the commercial mainframe industry.

4.5 John McCarthy (1927–2011) — The Father of Artificial Intelligence

American computer scientist and cognitive scientist John McCarthy made monumental contributions that shaped programming languages, operating systems, and intelligent machines.

- **Coining the Term "Artificial Intelligence" (1955):** McCarthy coined the phrase *Artificial Intelligence* in his formal proposal for the historic 1956 **Dartmouth Summer Research Project on Artificial Intelligence**, establishing AI as an autonomous, formal academic discipline.
- **Creation of LISP (1958):** Invented **LISP (List Processing)**, the second-oldest high-level programming language in continuous use. LISP introduced revolutionary concepts: dynamic memory structures, symbolic computation, recursive functions, higher-order functions, and abstract syntax trees.
- **Invention of Automatic Garbage Collection (1959):** To manage dynamic memory allocation in LISP without manual pointer reclamation, McCarthy invented **Garbage Collection** — the memory safety mechanism now universal in Java, Python, JavaScript, and C#.
- **Time-Sharing Operating Systems:** Pioneered the concept of computer time-sharing at MIT (CTSS), advocating for utility computing where thousands of remote users could interact with a central computing pool simultaneously.

Unit 5: Computing Systems: Past to Present Architectures

5.1 Processor Micro-Architectural Evolution

Microprocessor design has undergone a radical architectural transformation over the past five decades, driven by semiconductor physics and the pursuit of extreme parallel throughput.

1. Single-Core Processors & Frequency Scaling Limits

Early microprocessors contained a single physical execution engine that executed instructions in strict serial order. Designers scaled single-core performance through:

- **Instruction Pipelining:** Dividing the instruction cycle into discrete hardware stages (Fetch, Decode, Execute, Memory Access, Write-Back), enabling multiple instructions to be processed simultaneously at different stages like an assembly line.
- **Superscalar Execution:** Deploying multiple parallel execution units (ALUs, FPUs) to dispatch and complete multiple instructions per clock cycle ($IPC > 1$).
- **Out-of-Order (OoO) Execution:** Dynamically reordering instructions based on operand availability to avoid pipeline stalls.
- **The Breakdown of Dennard Scaling & The Power Wall (circa 2005):** Under classical **Dennard Scaling**, as transistor dimensions shrank by a factor of k , capacitance dropped and operating voltage could be scaled down proportionally, keeping power density constant ($P = CV^2f$). However, below **90 nm**, quantum tunneling leakage currents through atomic-scale gate oxides exploded. Operating voltage could no longer be reduced safely without compromising transistor stability. CPU power dissipation and thermal heat flux increased exponentially, hitting the **Power Wall** at ~4 GHz clock speeds, halting the single-core frequency race.

2. Dual-Core and Multi-Core Processors (CMP)

To continue performance scaling within acceptable thermal budgets (Thermal Design Power — TDP), the industry transitioned from monolithic single-core designs to **Chip Multiprocessors (CMP)** — integrating two, four, eight, or dozens of independent processing cores onto a single silicon die.

- **Thread-Level Parallelism (TLP):** Multiple execution threads execute concurrently across distinct physical CPU cores.
- **Simultaneous Multithreading (SMT) / Hyper-Threading:** Duplicating architectural register state (PC, registers) on a single physical core, presenting two logical processors to the operating system to maximize execution engine utilization during memory stalls.
- **Hierarchical Cache Coherence:** In multi-core processors, each core maintains private low-latency L1 and L2 caches, while sharing a massive Last-Level Cache (L3 Cache). When multiple cores hold copies of the same shared main memory location, hardware **Cache Coherence Protocols (MESI — Modified,**

Exclusive, Shared, Invalid, and MOESI) snoop the on-chip ring bus or mesh network to invalidate or update stale cache lines automatically.




Figure 1.2: Multi-Core Chip Multiprocessor (CMP) Architecture with Multi-Level Cache Hierarchy.

5.2 Graphics Processing Units (GPUs) & SIMD/SIMT Architecture

While a CPU is a latency-optimized processor engineered with massive caches and branch predictors to execute complex, sequential control flow with minimal delay, a **Graphics Processing Unit (GPU)** is a throughput-optimized processor designed for massive data-parallel computational workloads.

- **Architectural Organization:** A modern GPU (e.g., NVIDIA Hopper/Blackwell, AMD RDNA) contains thousands of compact, power-efficient Arithmetic Logic Units (ALUs) grouped into hundreds of Streaming Multiprocessors (SMs).
- **Execution Model:** Operates on the **SIMD (Single Instruction, Multiple Data)** and **SIMT (Single Instruction, Multiple Threads)** models. A single instruction is broadcast and executed simultaneously across thousands of parallel threads operating on distinct vector/matrix data elements.
- **General-Purpose GPU (GPGPU) Computing:** The release of frameworks like NVIDIA CUDA and OpenCL transformed GPUs into the primary computational engines for Deep Learning, Large Language Model training, computer vision, molecular dynamics, and cryptography.

5.3 Accelerated Processing Units (APUs) & Unified Memory

An **Accelerated Processing Unit (APU)** (pioneered by AMD and conceptually mirrored in Apple Silicon M-Series Unified Memory SoCs) integrates high-performance multi-core CPU cores and a powerful GPU onto a single monolithic silicon die.

- **Heterogeneous Unified Memory Architecture (hUMA):** In conventional PC systems, data must be copied across the high-latency PCIe expansion bus between system RAM and GPU VRAM. An APU allows CPU and GPU cores to share a unified physical memory address space, eliminating costly memory copies and reducing latency.
- **Benefits:** Exceptional energy efficiency, compact physical form factors, and optimized graphics performance in laptops, handheld gaming systems (Steam Deck), and home consoles (PlayStation 5, Xbox Series X).

5.4 Quantum Processing Units (QPUs) — Fundamental Concepts

A **Quantum Processing Unit (QPU)** is a computational device that harnesses the bizarre principles of quantum mechanics to solve certain intractable mathematical problems exponentially faster than classical supercomputers.

Core Principles of Quantum Computation:

1. **The Qubit (Quantum Bit):** Unlike a classical bit which is deterministically either **0** or **1**, a qubit is a two-level quantum system represented using Dirac bra-ket notation as a linear combination (superposition):

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle$$

where $\alpha, \beta \in \mathbb{C}$ are complex probability amplitudes satisfying the normalization constraint:

$$|\alpha|^2 + |\beta|^2 = 1$$

Upon measurement, the quantum superposition collapses irreversibly to classical state $|0\rangle$ with probability $|\alpha|^2$ or to state $|1\rangle$ with probability $|\beta|^2$.

2. **Quantum Superposition:** An array of n entangled qubits can exist simultaneously in a superposition of all 2^n computational basis states:

$$|\Psi\rangle = \sum_{k=0}^{2^n-1} c_k |k\rangle$$

While n classical bits represent exactly *one* value at any instant, n qubits hold all 2^n combinations simultaneously (**300 qubits $\approx 2^{300}$ states**, vastly exceeding the number of atoms in the observable universe).

3. **Quantum Entanglement:** A uniquely quantum phenomenon where two or more qubits become intimately correlated such that the quantum state of one qubit cannot be described independently of the others, regardless of physical separation (e.g., Bell State $|\Phi^+\rangle = \frac{1}{\sqrt{2}}(|00\rangle + |11\rangle)$).
4. **Quantum Interference & Algorithms:** Quantum algorithms (such as **Shor's Algorithm** for polynomial-time integer factorization and **Grover's Algorithm** for $O(\sqrt{N})$ unstructured database search) utilize quantum logic gates (Hadamard H , Pauli-X, CNOT) to constructively amplify the probability amplitudes of correct solutions while destructively cancelling erroneous paths.

Classical Computers vs Quantum Processors (QPUs)

Feature / Metric	Classical Computer (CPU / GPU)	Quantum Processing Unit (QPU)
Fundamental Information Unit	Binary Bit (0 or 1)	Qubit ($\alpha 0\rangle + \beta 1\rangle$)
Governing Physical Laws	Classical Electrodynamics, CMOS voltage states	Quantum Mechanics (Superposition, Entanglement, Interference)
Operating Environment	Room temperature (20°C to 70°C)	Near absolute zero ($\approx 15\text{ mK}$) in cryogenic dilution refrigerators
Error Susceptibility	Virtually zero (10^{-17} soft error rate)	Extremely high susceptibility to environmental thermal decoherence
Optimal Problem Domains	General software, operating systems, databases, web	Molecular simulation, quantum chemistry, cryptography (RSA factoring)